

Natural Language Processing

Episode 10-ish

Model Compression & Acceleration




Institute of
Science and
Technology
Austria


by Andrei Panferov

Chapter 1: Why Should You Care?

Case Study: Text Classification



Case Study: Text Classification




Three kinds of efficiency:



Model Size
(mega)bytes

Case Study: Text Classification



Three kinds of efficiency:




Model Size
(mega)bytes



Throughput
samples/second

Case Study: Text Classification



Three kinds of efficiency:



Model **S**ize
(mega)bytes




Throughput
samples/second

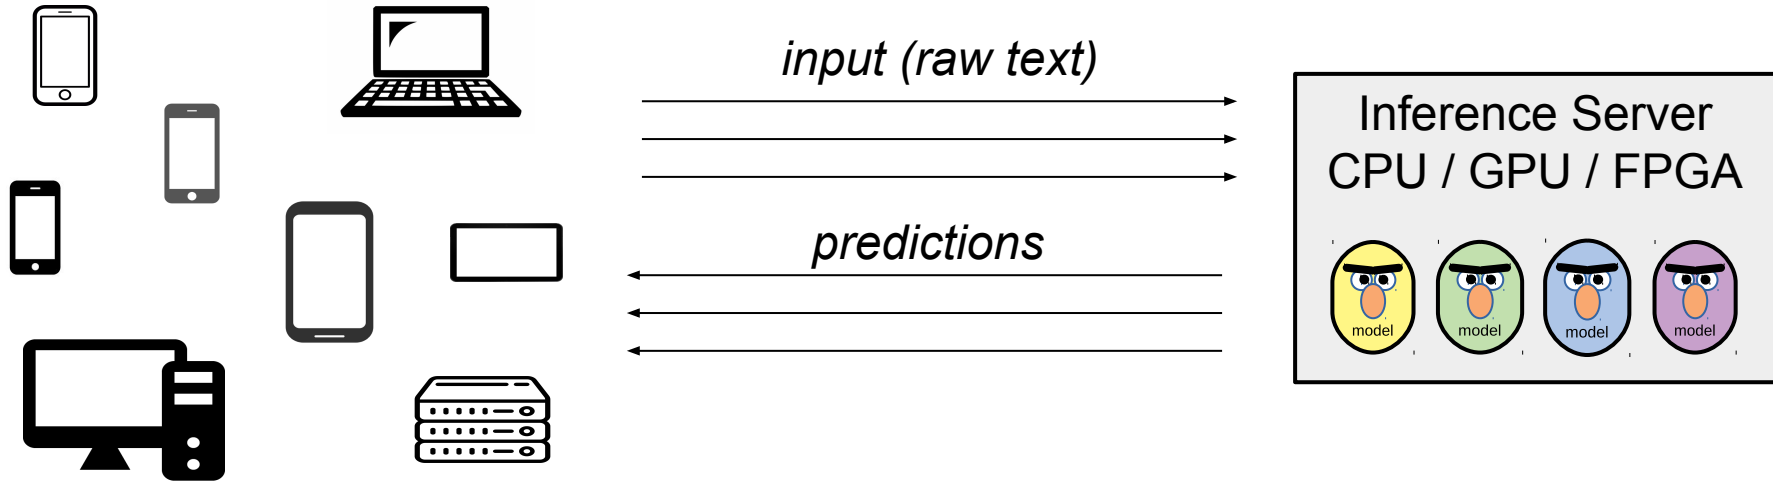


Latency
ms@percentile


Scenario 1: Inference Server



Scenario 1: Inference Server




Scenario 1: Inference Server



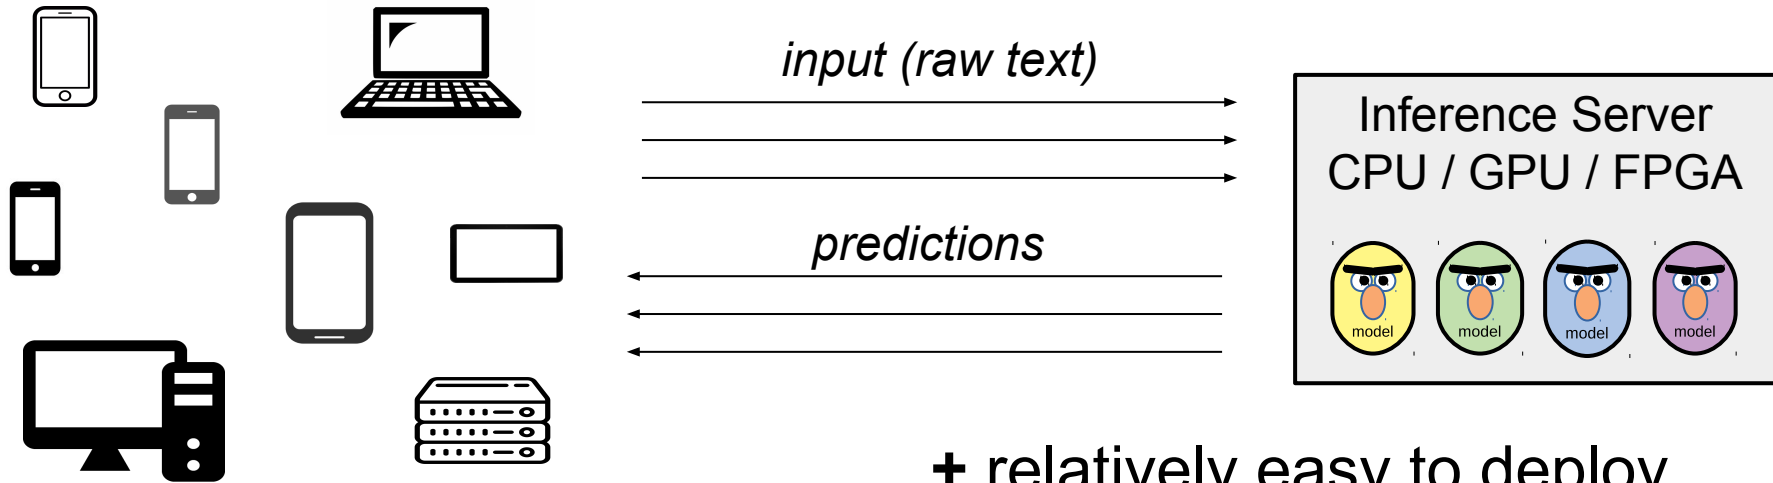
- + relatively easy to deploy
- + you control model & inference
- + clients don't run compute

Scenario 1: Inference Server



- + relatively easy to deploy
- + you control model & inference
- + clients don't run compute
- you pay for each inference
- clients can't work offline
- network latency

Scenario 1: Inference Server

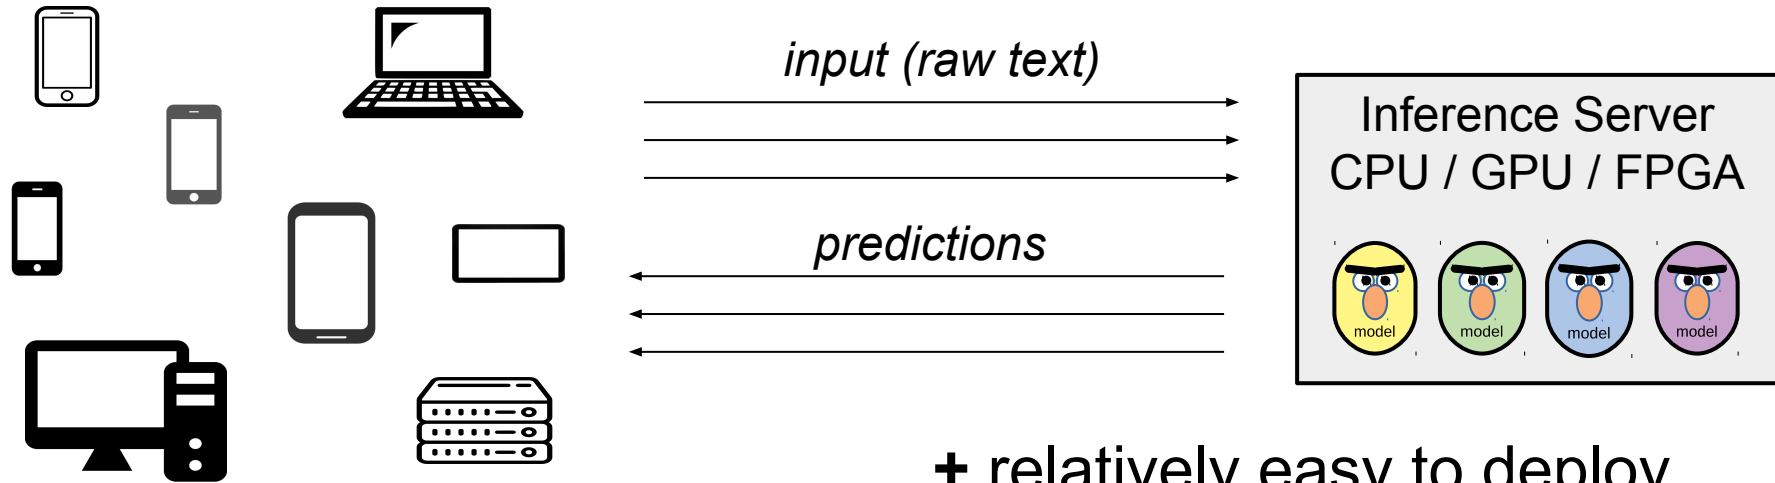


Which is the most important?



- + relatively easy to deploy
- + you control model & inference
- + clients don't run compute
- you pay for each inference
- clients can't work offline
- network latency

Scenario 1: Inference Server



Priorities:



- + relatively easy to deploy
- + you control model & inference
- + clients don't run compute
- you pay for each inference
- clients can't work offline
- network latency

Scenario 1: Inference Server

- Group inputs into batches (e.g. by length)
 - *improves throughput at the cost of latency*
- Multiple servers with load balancing

Scenario 1: Inference Server

- Group inputs into batches (e.g. by length)
 - *improves throughput at the cost of latency*
- Multiple servers with load balancing
 - improves throughput at the cost of your budget :)*

Popular frameworks:



[vLLM](#)



[Triton Inference Server](#)



Custom model-dependent code

priorities

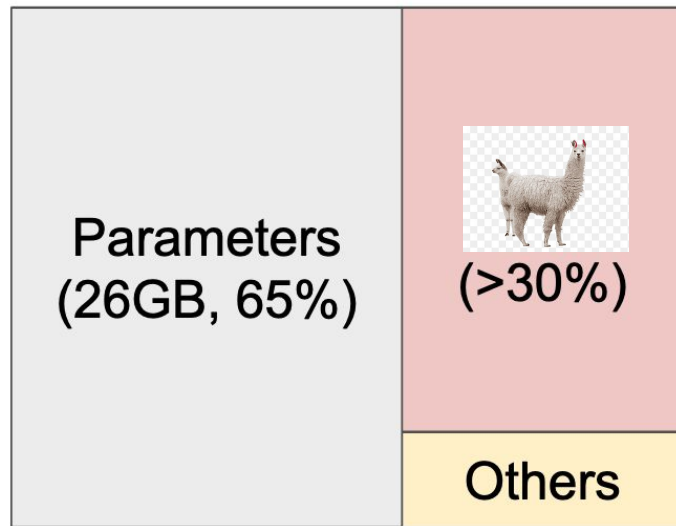
efficiency \ll developer time

efficiency \approx developer time

efficiency \gg developer time

Scenario 1: Inference Server

Question: what did behind the Llama?




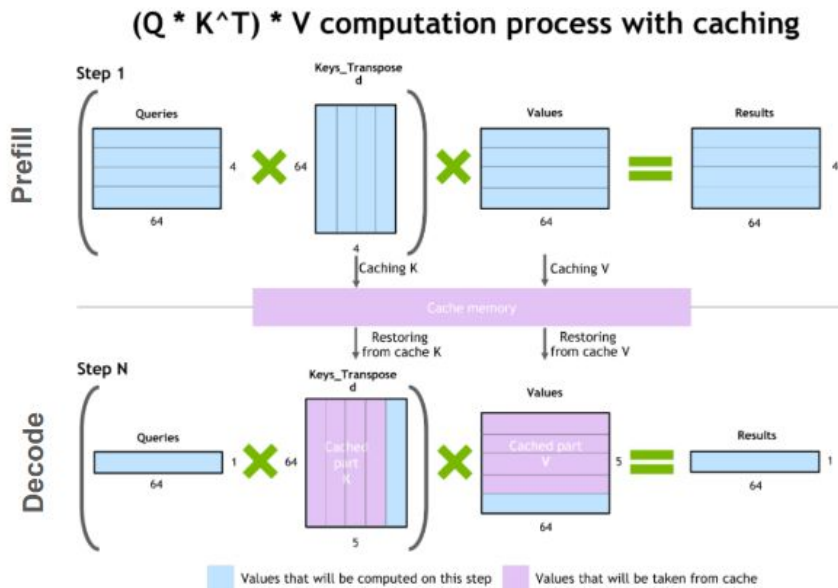
NVIDIA A100 40GB

Image: arxiv.org/pdf/2309.06180

Scenario 1: Inference Server

Question: what did behind the Llama?

Answer: KV Cache




NVIDIA A100 40GB

Image: arxiv.org/pdf/2309.06180

Scenario 1: Inference Server

Question: how many parallel requests we can serve if we have 20Gb of extra DRAM and one token takes up 200kb of cache?

Answer: it depends...



NVIDIA A100 40GB

Image: arxiv.org/pdf/2309.06180


Scenario 1: Inference Server

Question: how many parallel requests we can serve if we have 20Gb of extra DRAM and one token takes up 200kb of cache?

Answer: it depends...

What if requests were very diverse:

- 2000 prompt tok, 1 gen tok
- 2000 prompt tok, 2000 gen tok
- 50 prompt tok, 1 gen tok



NVIDIA A100 40GB

Image: arxiv.org/pdf/2309.06180


Scenario 1: Inference Server

Question: how many parallel requests we can serve if we have 20Gb of extra DRAM and one token takes up 200kb of cache?

Answer: it depends...

What if requests were very diverse:

- 2000 prompt tok, 1 gen tok
- 2000 prompt tok, 2000 gen tok
- 50 prompt tok, 1 gen tok



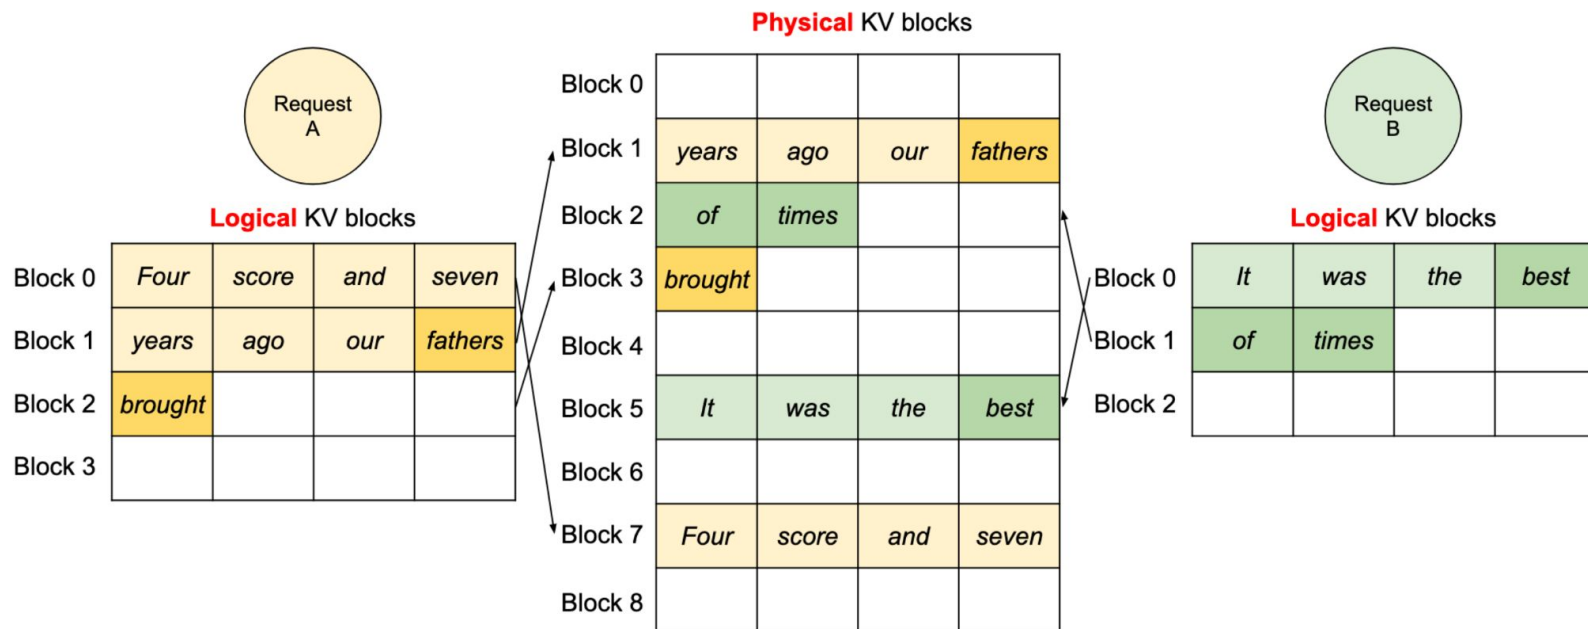
NVIDIA A100 40GB

Image: arxiv.org/pdf/2309.06180

Scenario 1: Inference Server

Question: how many parallel requests we can serve if we have 20Gb of extra DRAM and one token takes up 200kb of cache?

Answer: ~100000 tokens with Paged Attention



Scenario 1: Inference Server

TLDR:

- Prioritizing throughput
- Many concurrent users with diverse requests
- Smart KV-Cache management is a must

Scenario 2: Workstation Deployment

- Preload model onto a dedicated device, infer locally
- Features:
 - Data privacy
- Requirements:
 - Predictable hardware
- LLM frameworks:
 - llama.cpp
 - ollama
 - exllama-v2

Priorities:



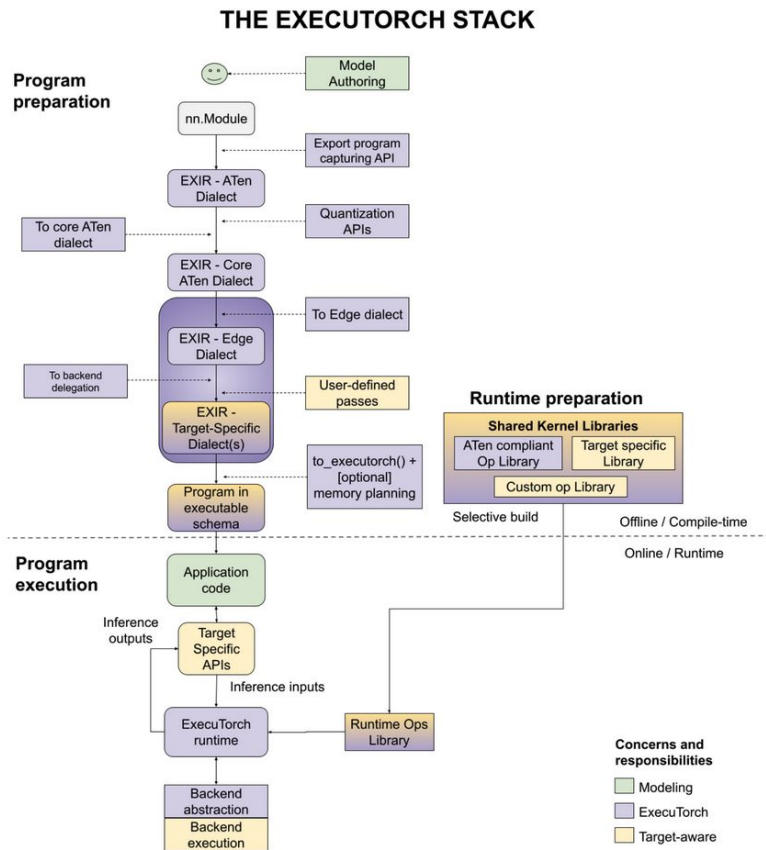
Scenario 3: Mobile Deployment

- Run model on a mobile device
- Features:
 - Diverse devices (or not if you're Apple)
 - New priority: *power consumption*
 - Specialized matmul kernels

Priorities:



Scenario 3: Mobile Deployment



Same model...

- Tracing the model
- Compiling

Different backends

- Mapping to kernels

Scenario 3: Mobile Deployment

 [TensorFlow.js](#)

 [CoreML](#)

 [NNAPI](#)

 [ExecuTorch](#)

 [MLX](#)

All modern browsers

iOS devices

Android devices

iOS/Android/Embedded

Apple Silicon

Scenario 3: Mobile Deployment

TLDR:

- Often limited by RAM/Latency
- Sensitive to power consumption
- Specialized/unique/diverse hardware


Chapter 2: How Do I Improve My Model?

Compression by Distillation



Distillation...
Heard that word before?

Compression by Distillation

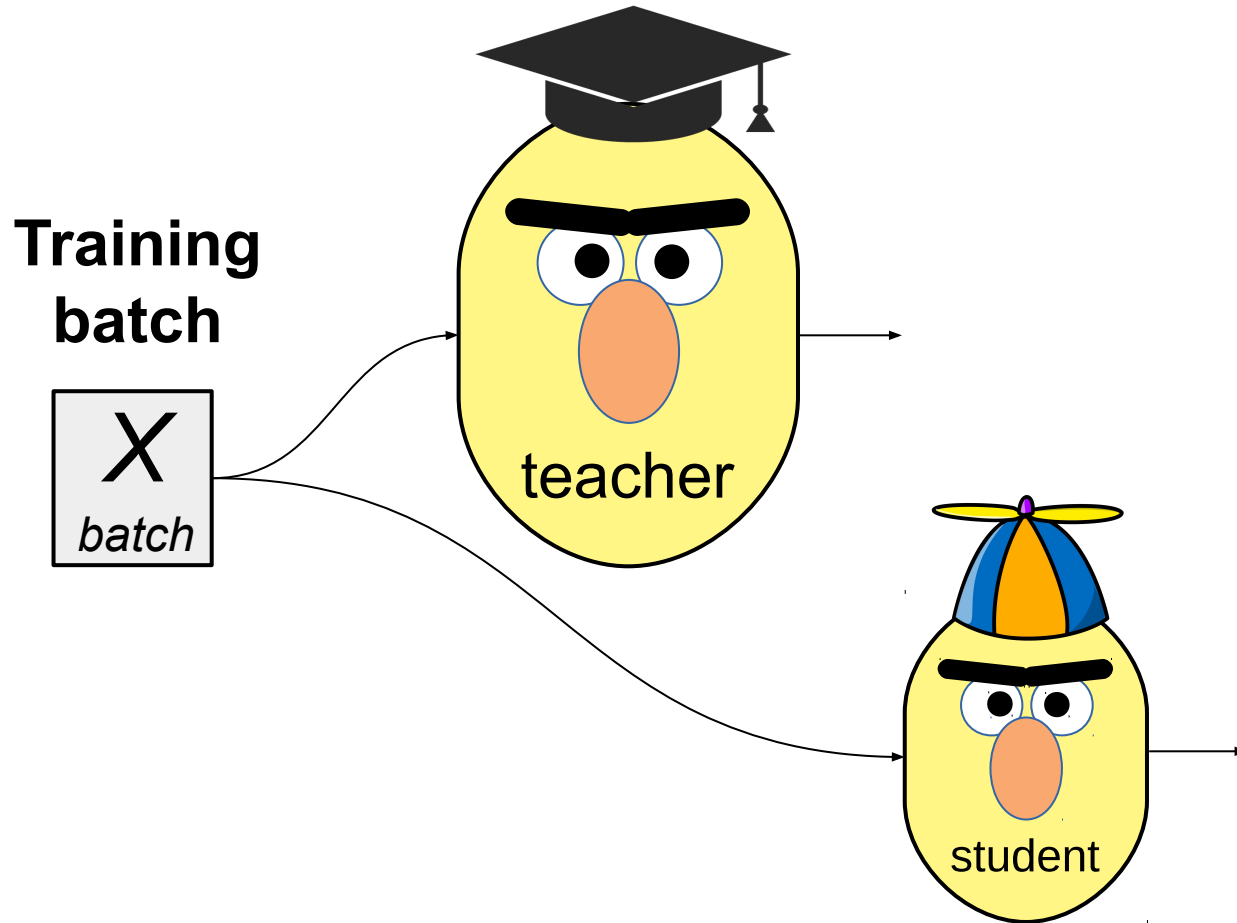


First, get the best performing model regardless of size

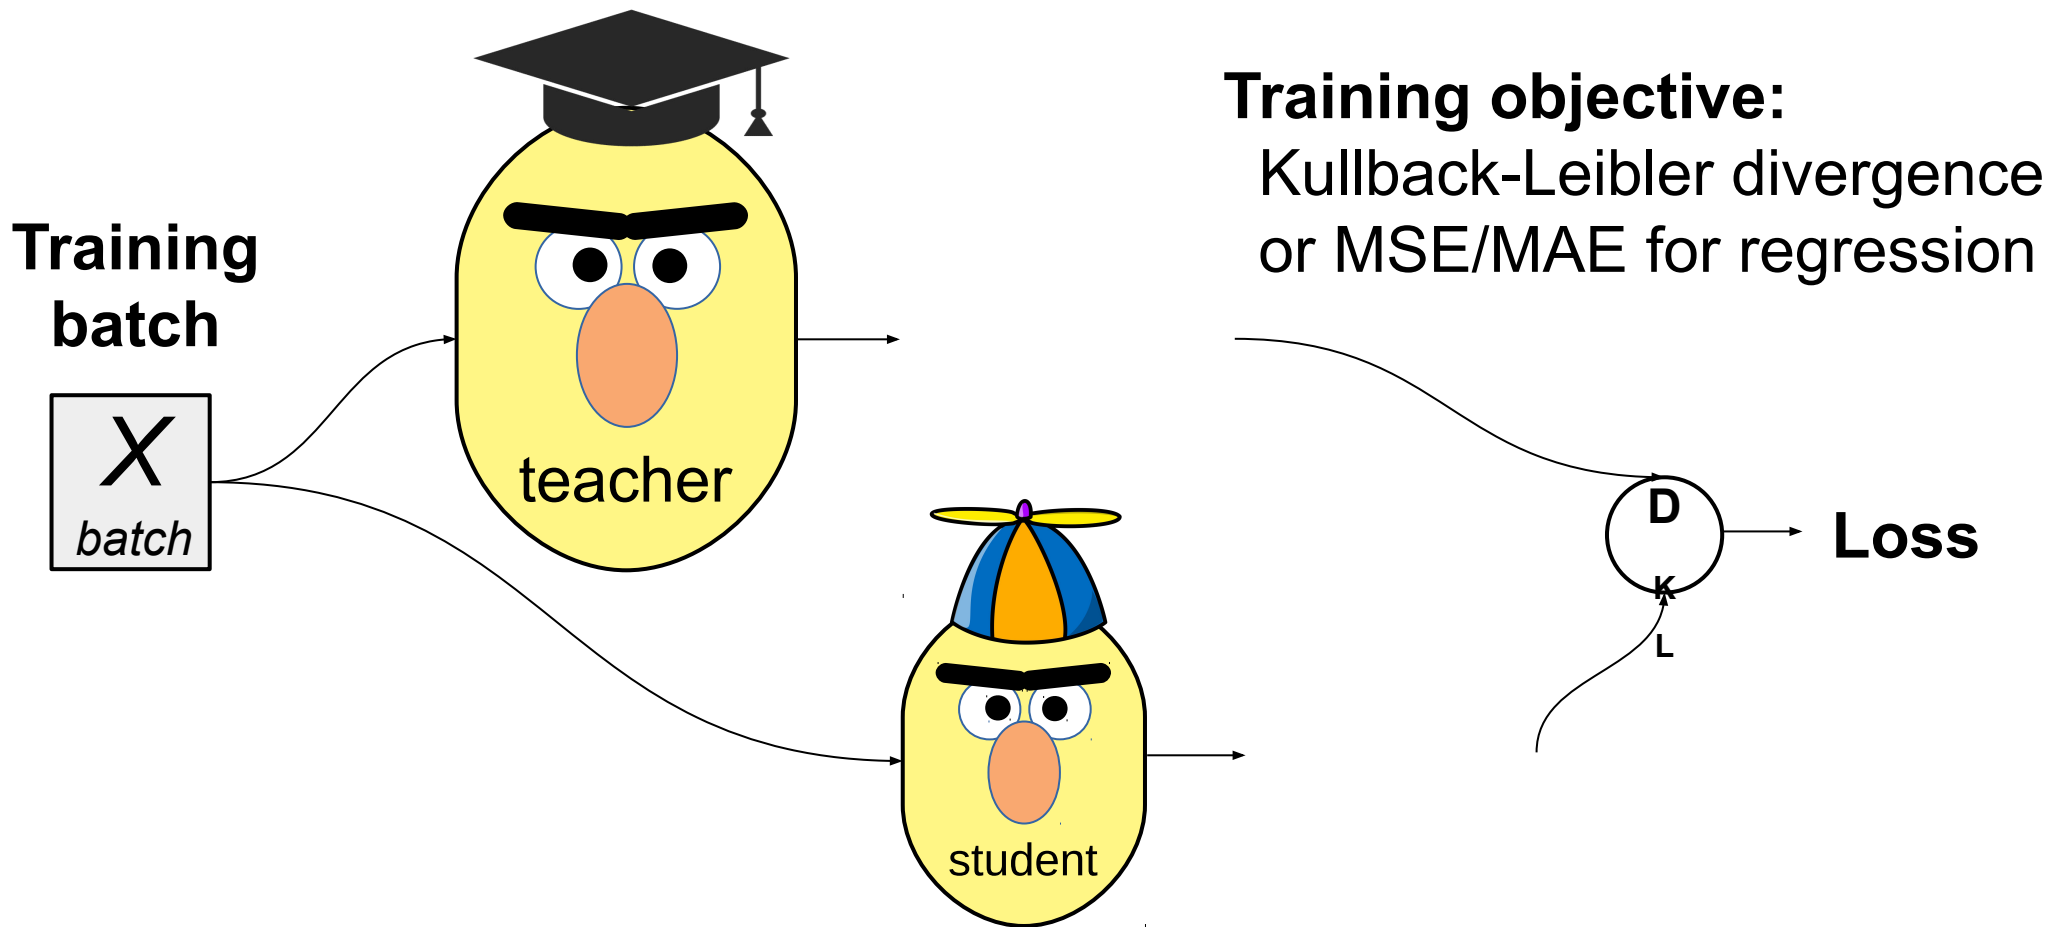


Then, train a more compact model to approximate it!

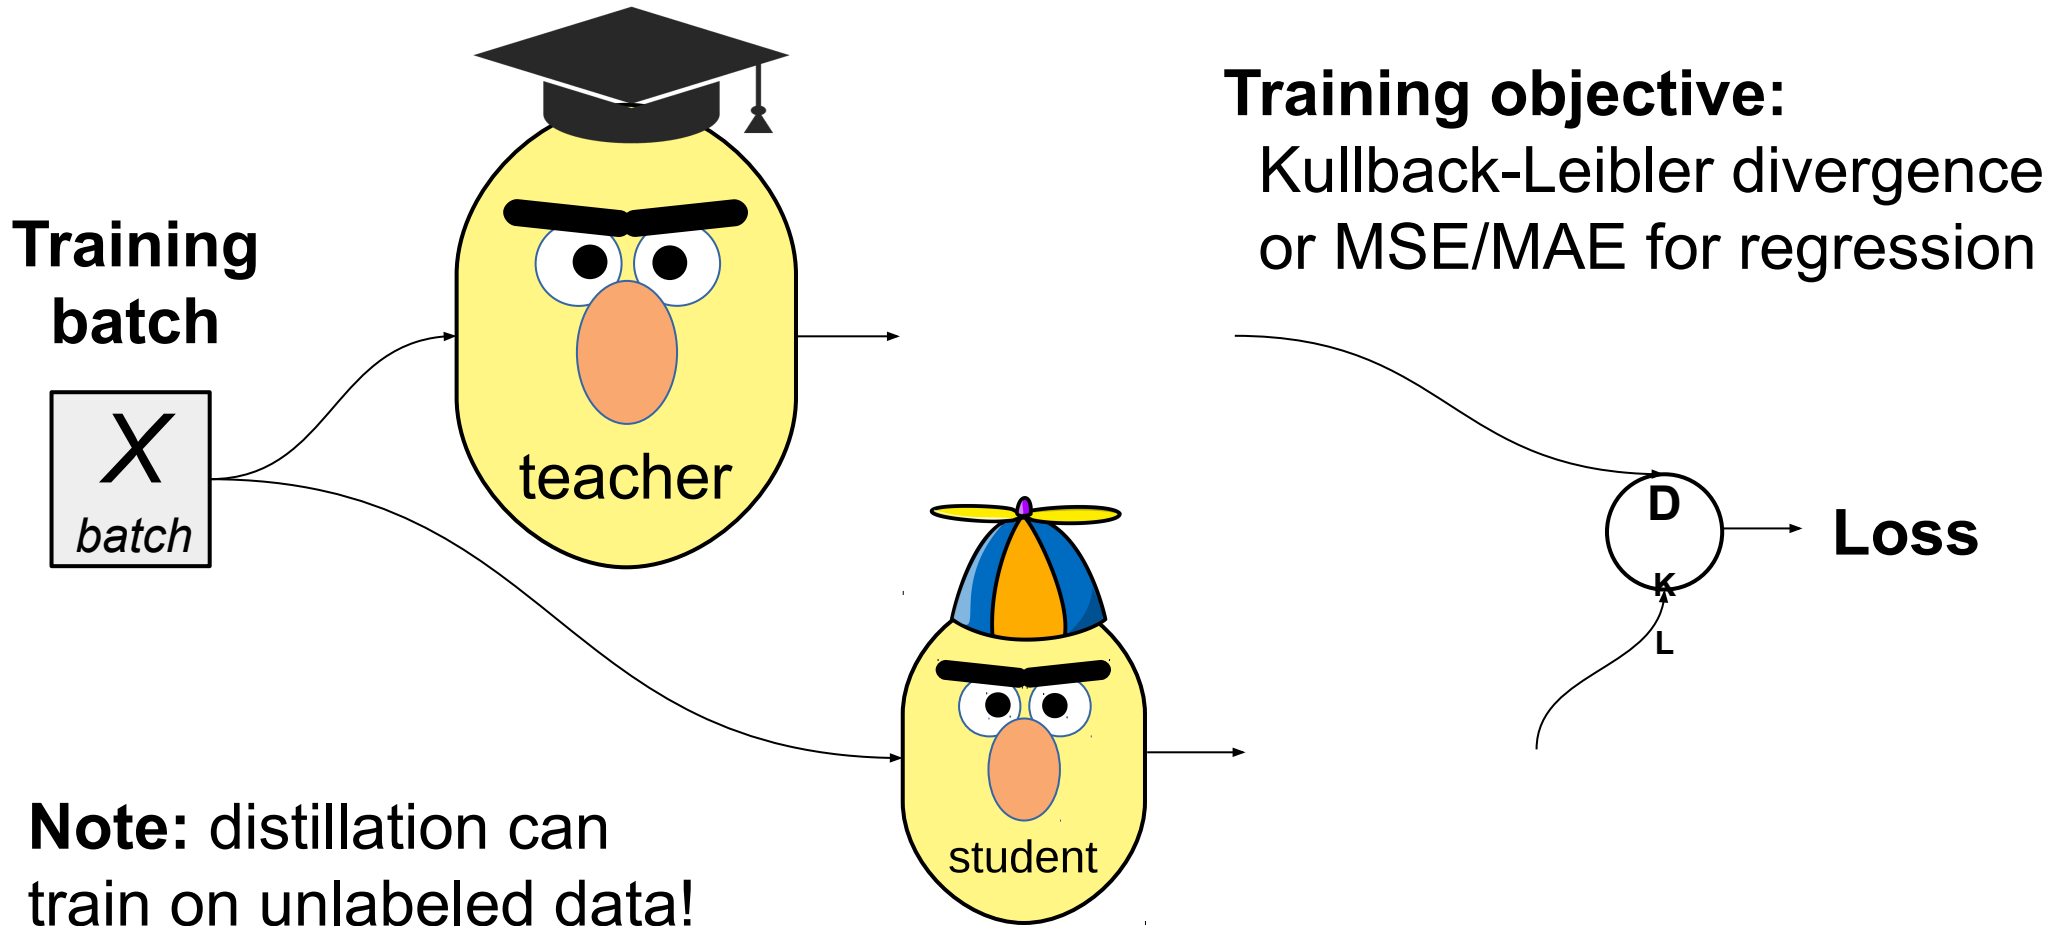
Compression by Distillation



Compression by Distillation



Compression by Distillation



Compression by Distillation

- Student architecture choices:

Naïve: same but smaller, less layers / hidden units

- e.g. DistilBERT: <https://arxiv.org/pdf/1910.01108.pdf>
- Same as BERT-base, but with
- half as many layers*
- (and ≈ 1.5 times faster)

Model	# param. (Millions)	Inf. time (seconds)
ELMo	180	895
BERT-base	110	668
DistilBERT	66	410

Model	Score	CoLA	MNLI	MRPC	QNLI	QQP	RTE	SST-2	STS-B	WNLI
ELMo	68.7	44.1	68.6	76.6	71.1	86.2	53.4	91.5	70.4	56.3
BERT-base	79.5	56.3	86.7	88.6	91.8	89.6	69.3	92.7	89.0	53.5
DistilBERT	77.0	51.3	82.2	87.5	89.2	88.5	59.9	91.3	86.9	56.3

Compression by Distillation

- Student architecture choices:
 - Naïve:** same but smaller, less layers / hidden units, random init
 - **Sparse along layers/attn_heads/mlp_neurons:** sparsify but obtain dense representation, init from teacher
 - Let's focus on that...

Minitron Approach

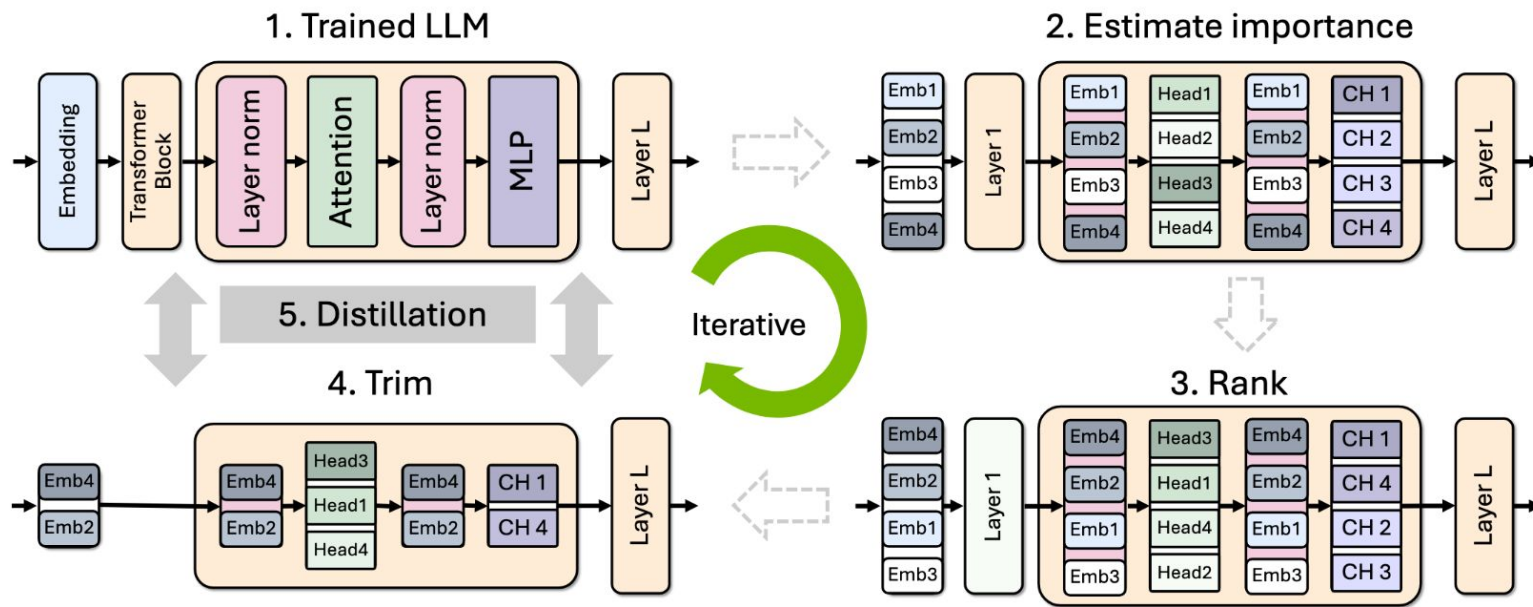



Image: arxiv.org/pdf/2407.14679

Figure 2: High-level overview of our proposed iterative pruning and distillation approach to train a family of smaller LLMs. On a pretrained LLM, we first evaluate importance of neurons, rank them, trim the least important neurons and distill the knowledge from the original LLM to the pruned model. The original model is replaced with the distilled model for the next iteration of compression.

Minitron Approach



Source: arxiv.org/pdf/2407.14679

Minitron Approach

1. To train a family of LLMs, train the largest one and prune+distill iteratively to smaller LLMs.
2. Use (batch=L2, seq=mean) importance estimation for width axes and PPL/BI for depth.
3. Use single-shot importance estimation; iterative provides no benefit.
4. Prefer width pruning over depth for the model scales we consider ($\leq 15B$).
5. Retrain exclusively with distillation loss using KLD instead of conventional training.
6. Use (logit+intermediate state+embedding) distillation when depth is reduced significantly.
7. Use logit-only distillation when depth isn't reduced significantly.
8. Prune a model closest to the target size.
9. Perform lightweight retraining to stabilize the rankings of searched pruned candidates.
10. If the largest model is trained using a multi-phase training strategy, it is best to prune and retrain the model obtained from the final stage of training.

Compression by Distillation

- More distillation tricks:

- Ensemble distillation

<https://arxiv.org/abs/1702.01802>


- Dropout distillation

<http://proceedings.mlr.press/v48/bulo16.pdf>


- Co-distillation

<https://arxiv.org/abs/1804.03235>

Smaller Model for Simpler Inputs



Smaller LLMs for Simpler Tokens?



Bottlenecks in GPU Inference

The math behind GPU utilization.

K tokens at a time:

- **Memory transfer** load of $O(N^2 + K \times N)$ per matrix-vector product.
- **Compute** load of $O(K \times N^2)$ per matrix-vector product.

Time cost of GPU operations:

- Bring million numbers from memory into kernels - **SLOW**
- Multiply million numbers with kernels - **FAST**

Bottlenecks in GPU Inference

$K \approx 1$ token at a time:

- The primary load in single user **chat applications**.
- **Memory transfer** load of $O(N^2)$ per matrix-vector product.
- **Compute** load of $O(K \times N^2)$ per matrix-vector product.
- Memory transfer bottlenecked.

$K > N$ tokens at a time:

- Prompt preprocessing, **highly parallel** inference.
- **Memory transfer** load of $O(K \times N)$ per matrix-matrix product.
- **Compute** load of $O(K \times N^2)$ per matrix-matrix product.
- Compute bottlenecked.

Bottlenecks in GPU Inference

$K \approx 1$ token at a time:

- Latency doesn't depend on K !

$K > N$ tokens at a time:

- Latency is linear with K .


Speculative Decoding

<https://arxiv.org/abs/2211.17192>

Two models: main model (~Llama-70B) and draft model (~Llama-7B)

Greedy decoding:

Step 1: generate with draft model (sequential)



Speculative Decoding



<https://arxiv.org/abs/2211.17192>

Two models: main model (~Llama-70B) and draft model (~Llama-7B)

Greedy decoding:

Step 1: generate with draft model (sequential)

Step 2: verify with large model (parallel)



Speculative Decoding

<https://arxiv.org/abs/2211.17192>


Two models: main model (~Llama-70B) and draft model (~Llama-7B)

Greedy decoding:

Step 1: generate with draft model (sequential)

Step 2: verify with large model (parallel)

Accept multiple tokens



Speculative Decoding

<https://arxiv.org/abs/2211.17192>

Two models: main model (~Llama-70B) and draft model (~Llama-7B)


Greedy decoding:

Step 1: generate with draft model (sequential)

Step 2: verify with large model (parallel)

Accept multiple tokens

Repeat

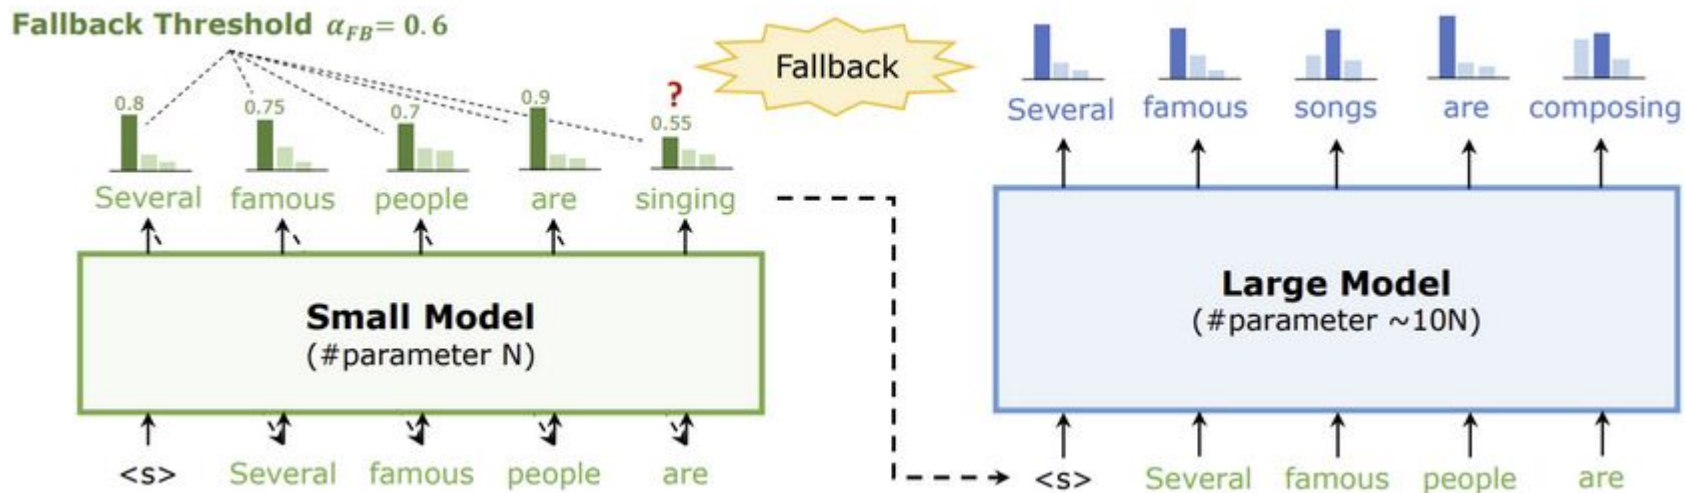


Speculative Decoding

<https://arxiv.org/abs/2211.17192>

Two models: main model (~Llama-70B) and draft model (~Llama-7B)

Sampling (temperature, top-p, top-k): generate, then reject with probability sampling probability proven equal to regular sampling




Speculative Decoding

That you'll implement

Two models: main model (Llama-8B) and draft model (of your choice)

Greedy Sampling: generate, then reject everything after exact match




Going Inside Model Layers

Compression by Sparsification

Do we really need all D by D weights?

Compression by Pruning

Do we really need all D by D weights?



Magnitude Pruning

Drop ~5% smallest weights
from each layer every 1000 steps
(and keep training)

Importance Estimation

Minitron:

- attn_head/neuron level: activation scales
- Block level: PPL

L0: Learnable importance estimation:

$$w_i = w_i \times \sigma(a_i + N(0, 1))$$

Read more: <https://arxiv.org/abs/1712.01312>

Alternative: <https://arxiv.org/abs/1701.05369>

Compression by Sparsification

Unstructured sparsity = prune individual weights
(minimal model size)


Structured sparsity= prune entire neurons/heads
(fastest inference)

Compression by Sparsification





Unstructured sparsity = prune individual weights
(minimal model size)

Structured sparsity= prune entire neurons/heads
(fastest inference)

Note: some GPU/FPGAs also run fast
with low-level structured sparsity, e.g.
“Any 2 of 4 consecutive weights” (left)



Compression by quantization




INT8

8 BITS


Quantization basics

Goal: encode data as int8 / int4



Quantization basics


Goal: encode data as int8 / int4



not an ideal range for int4

Linear quantization

Fit a linear range to data



$$\text{scale} = (\max(w) - \min(w)) / 2^4$$


$$\text{zero} = -\min(w) / \text{scale}$$

Linear quantization

Fit a linear range to data

Encode: $\mathbf{c}_i = (\mathbf{w}_i / \mathbf{s} + \mathbf{z}).\text{clip}(0, 15)$ uint4 range

Decode: $\mathbf{w}_i = ???$ ideas?




$$\mathbf{scale} = (\max(\mathbf{w}) - \min(\mathbf{w})) / 2^4$$
$$\mathbf{zero} = -\min(\mathbf{w}) / \mathbf{scale}$$

Linear quantization

Fit a linear range to data

Encode: $\mathbf{c}_i = (\mathbf{w}_i / \mathbf{s} + \mathbf{z}).\text{clip}(0, 15)$ uint4 range

Decode: $\mathbf{w}_i \approx \mathbf{s} * \mathbf{c}_i - \mathbf{z}$



$\mathbf{scale} = (\max(\mathbf{w}) - \min(\mathbf{w})) / 2^4$
 $\mathbf{zero} = -\min(\mathbf{w}) / \mathbf{scale}$


LLM.int8(): some weights are more important

<https://arxiv.org/abs/2208.07339>

TL;DR in very LLM, some input features become outliers


Weights for those features are sensitive

KEEP <1% MOST SENSITIVE WEIGHTS IN 16-bit!



Nonlinear quantization

Fit codes to data




Compute a grid of percentiles or centroids (k-means 1d)

Store each weight as the index of nearest percentile/centroid

Static nonlinear case: NF4

<https://arxiv.org/abs/2305.14314>

<https://arxiv.org/abs/2306.06965>



Static nonlinear case: NF4

<https://arxiv.org/abs/2305.14314>

<https://arxiv.org/abs/2306.06965>

How to use:

```
1 model = transformers.AutoModelForCausalLM.from_pretrained(  
2     "Enoch/llama-7b-hf", load_in_4bit=True)
```

`[-1. , -0.7 , -0.53, -0.39, -0.28, -0.18, -0.09, 0. , 0.08, 0.16, 0.25, 0.34, 0.44, 0.56, 0.72, 1.]`



7 negative quantiles zero (exact) 8 positive quantiles

How many bits is best?

<https://arxiv.org/abs/2212.09720>

TL;DR 3-4 bits looks optimal

2-bit: Smaller models in 4 bits are better than larger models in 2 bits




How many bits is best?

<https://arxiv.org/abs/2401.06118>

TL;DR ~2 bits looks optimal

2-bit: more expensive to quantize, more expensive to run



GPU Inference

$K \approx 1$ token at a time:

- The primary load in single user **chat applications**.
- **Memory transfer** load of $O(N^2)$ per matrix-vector product.
- **Compute** load of $O(K \times N^2)$ per matrix-vector product.
- Memory transfer bottlenecked.

$K > N$ tokens at a time:

- Prompt preprocessing, **highly parallel** inference.
- **Memory transfer** load of $O(K \times N)$ per matrix-matrix product.
- **Compute** load of $O(K \times N^2)$ per matrix-matrix product.
- Compute bottlenecked.

Quantized GPU Inference

With weight compression ratio **C**.

$K \approx 1$ token at a time:

- The primary load in single user **chat applications**.
- **Memory transfer** load of $O(N^2 / C)$ per matrix-vector product.
- **Compute** load of $O(K \times N^2)$ per matrix-vector product.
- Memory transfer bottlenecked.

$K > N$ tokens at a time:

- Prompt preprocessing, **highly parallel** inference.
- **Memory transfer** load of $O(K \times N)$ per matrix-matrix product.
- **Compute** load of $O(K \times N^2)$ per matrix-matrix product.
- Compute bottlenecked.

Quantized GPU Inference

With weight compression ratio C .

$K \approx 1$ token at a time:

- **Memory transfer** load of $O(N^2 / C)$ per matrix-vector product.
- Memory transfer bottlenecked.

$K > N$ tokens at a time:

- **Compute** load of $O(K \times N^2)$ per matrix-matrix product.
- Compute bottlenecked.

Quantized GPU Inference

With weight compression ratio C .

$K \approx 1$ token at a time:

- C times faster inference!

$K > N$ tokens at a time:

- No speedup :(

Fully Quantized GPU Inference

With weight+activations compression ratio **C**.

$K \approx 1$ token at a time:

- **Memory transfer** load of $O(N^2 / C)$ per matrix-vector product.
- Memory transfer bottlenecked.

$K > N$ tokens at a time:

- **Compute** load of $O(K \times N^2 / C)$ per matrix-matrix product.
- Compute bottlenecked.

Fully Quantized GPU Inference

With weight+activations compression ratio **C**.

$K \approx 1$ token at a time:

- **C** times faster inference!

$K > N$ tokens at a time:

- **C** times faster inference!

Model compression landscape

Goal: faster / smaller / both

Compression: quantize / prune / factorize

Setup: no data, some data, training data

Compress what: weights / activations / cache

Model compression landscape

Goal: faster / smaller / both

Compression: quantize / prune / factorize

Setup: no data, some data, training data

Compress what: weights / activations / cache

Q: can we take advantage of data
to improve quantization?

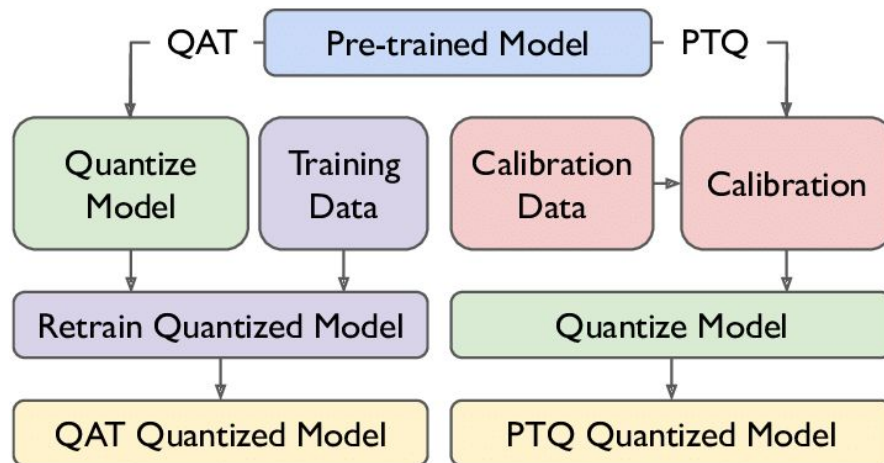
Model compression landscape

Goal: faster / smaller / both

Compression: quantize / prune / factorize

Setup: no data, some data, training data

Compress what: weights / activations / cache



Compression-aware training

Step 1: train normally for T steps

Step 2: prune 5% weights (or quantize 10% layers)

Step 3: freeze pruned/quantized parts

GoTo **step 1**

Post-training compression

Can we take advantage of **a few** data points without training the model?

Post-training compression

Can we take advantage of **a few** data points without training the model?

- find which weights are multiplied by larger inputs
- find correlated or anti-correlated weights
- try to “cancel out” quantization errors

GPTQ

[Frantar et al, 2022](#)


Minimize the same objective

$$\operatorname{argmin}_{\widehat{\mathbf{W}}_\ell} \|\mathbf{W}_\ell \mathbf{X}_\ell - \widehat{\mathbf{W}}_\ell \mathbf{X}_\ell\|_2^2$$

For $i = 1 \dots \text{in_features}$:


- **quantize** i -th column of weight matrix
(from one input feature and all outputs)
- freeze the quantized model forever
- update all remaining columns

Inverse Layer Hessian
(Cholesky Form)



computed initially

Weight Matrix / Block



block i quantized recursively
column-by-column



GPTQ

[Frantar et al, 2022](#)

Minimize the same objective

$$\operatorname{argmin}_{\widehat{\mathbf{W}}_\ell} ||\mathbf{W}_\ell \mathbf{X}_\ell - \widehat{\mathbf{W}}_\ell \mathbf{X}_\ell||_2^2$$


For $i = 1 \dots \text{in_features}$:

- **quantize** i -th column of weight matrix
(from one input feature and all outputs)
- freeze the quantized model forever
- update all remaining columns

Use linear quantization with one scale & zero per each group of G weights


Tricks: process weights in “hacky” order

Inverse Layer Hessian
(Cholesky Form)




computed initially

Weight Matrix / Block



block i quantized recursively
column-by-column



GPTQ

[Frantar et al, 2022](#)

Minimize the same objective

$$\operatorname{argmin}_{\widehat{\mathbf{W}}_\ell} ||\mathbf{W}_\ell \mathbf{X}_\ell - \widehat{\mathbf{W}}_\ell \mathbf{X}_\ell||_2^2$$

For i = 1

- **quantize**

(from

- freeze t

- update

Use linear

zero per e


Native support in HF Transformers:

<https://huggingface.co/blog/gptq-integration>

Original implementation:

<https://github.com/IST-DASLab/gptq>

Weight Matrix / Block



i quantized recursively
column-by-column

quantized weights
that are updated

Tricks: process weights in “hacky” order

SparseGPT

[Frantar et al, 2023](#)

Minimize the same objective

$$\operatorname{argmin}_{\widehat{\mathbf{W}}_{\ell}} ||\mathbf{W}_{\ell}\mathbf{X}_{\ell} - \widehat{\mathbf{W}}_{\ell}\mathbf{X}_{\ell}||_2^2$$


For $i = 1 \dots \text{in_features}$:

- **sparsify** i -th column of weight matrix
(from one input feature and all outputs)
- freeze the quantized model forever
- update all remaining columns

Dynamically choose how many weights to prune on each step (threshold on error)

Trick: can do 2-out-of-4 sparsity for A100

OPT-175B



QuIP#/AQLM/SpinQuant

Minimize the same objective $\operatorname{argmin}_{\widehat{\mathbf{W}}_\ell} ||\mathbf{W}_\ell \mathbf{X}_\ell - \widehat{\mathbf{W}}_\ell \mathbf{X}_\ell||_2^2$

QuIP#/AQLM/SpinQuant

Minimize the same objective $\operatorname{argmin}_{\widehat{\mathbf{W}}_\ell} ||\mathbf{W}_\ell \mathbf{X}_\ell - \widehat{\mathbf{W}}_\ell \mathbf{X}_\ell||_2^2$

But also finetune while quantizing!

QuIP#

Minimize the same objective $\operatorname{argmin}_{\widehat{\mathbf{W}}_\ell} \|\mathbf{W}_\ell \mathbf{X}_\ell - \widehat{\mathbf{W}}_\ell \mathbf{X}_\ell\|_2^2$

But also finetune while quantizing!

On the level of blocks, after quantizing each linear layer:

$$\operatorname{argmin}_{\text{Unquantized} \in B} \|B(X) - \widehat{B}(X)\|_2^2$$

More Quantization Papers

PV-Tuning – SOTA weight-only quantization optimization procedure

<https://arxiv.org/abs/2405.14852>

QTIP – SOTA weight-only quantization algorithm

<https://arxiv.org/abs/2406.11235>

QuaRot - SOTA GPTQ-based weight+act quantization algorithm

<https://arxiv.org/abs/2404.00456>

SpinQuant – SOTA weight+act quantization algorithm with finetuning

<https://arxiv.org/abs/2405.16406>

Many more cool papers

Homework:

- ``hw_speculative.ipynb`` - Implement simple speculative decoding
 - 4+(?) Points
- ``hw_quantization.ipynb`` - Implement GPTQ
 - 6+(3) Points